Beyond Detection: FARABI and the Assessment Credibility Shock in Higher Education

Ewan Simpson

Narxoz Business School, Narxoz University

January 8, 2026

Abstract

Generative AI is creating a credibility shock for higher education assessment. The issue is not only misconduct but a weakened evidentiary link between student work and the learning claims behind grades and credentials. Surveys and guidance from 2024 to 2025 indicate rapid uptake and that detection cannot provide assurance. This paper argues that traditional assessment systems, including take-home coursework graded mainly on the final product and invigilated exams, were already noisy evidence for higher-order learning. Generative AI accelerates these weaknesses by making plausible outputs cheap and by displacing tool influence into preparation and rehearsal. In response, the paper introduces FARABI (Framework for AI-Resilient Assessment and Balanced Integrity), a risk-based triage method for assessment portfolios. FARABI supports proportionate controls while acting as a diagnostic that shifts assessment toward applied judgment, observable process, and accountable tool use. The paper closes by questioning higher education's warrant if employability weakens as the organizing metric.

Keywords: artificial intelligence; higher education; AI governance; assessment integrity

1. Introduction: A known problem with weak response capacity

Higher education is facing an assessment credibility shock. In the hands of a skilled user, generative AI systems can now produce fluent text, plausible analyses, and polished presentations at scale, high speed and at low cost. The immediate institutional response has often centered on academic integrity rules, prohibited-use statements, and detection. Yet the deeper problem is evidentiary. When a graded artifact can be manufactured quickly with tools that are hard to observe and harder to attribute, the artifact becomes weaker evidence of what the student knows or can do. At that point, confidence in grades and credentials is no longer a matter of isolated misconduct. It becomes a system-level credibility issue.

Two features make this problem structurally different from earlier integrity challenges. First, adoption is fast. Recent survey evidence in the UK shows steep year-on-year jumps in student use, including use for assessment tasks (Freeman, 2025). Global snapshots similarly indicate widespread student use and expectations that universities provide clear guidance and support for appropriate use (DEC, 2024). Second, the tools are general-purpose. They can assist with ideation, structure, drafting, translation, code, and synthesis across many disciplines. That combination of rapid diffusion plus broad capability turns assessment into a moving target and reduces the viability of incremental fixes.

At the same time, this is not a clean break from an otherwise robust baseline. The credibility shock lands on an assessment ecosystem that has long been contested on validity, signaling, and integrity. Many institutions relied heavily on a familiar set of take-home assignments, essays, reports, problem sets, and presentations completed outside controlled conditions. These formats persist partly because they are administratively convenient and scale well. They also align with a widely held academic identity: students learn by producing written products and educators judge those products. However, that model is often a weak way to evidence higher-order learning at scale, even without AI.

1.1 The traditional coursework model was already weak evidence

This paper uses the term 'traditional coursework model' to mean take-home assignments graded mainly on the final product, with limited visibility into process. The core claim is not that such tasks are worthless. It is that, as evidence for higher-order learning, they have long been vulnerable and noisy.

From a validity perspective, assessment is an argument about interpretation and use: what a score is taken to mean, and what decisions it justifies (Kane, 2013; Messick, 1995). When the primary evidence is a polished product created outside secure conditions, alternative explanations for performance multiply. The problem is not only plagiarism in the narrow sense, but the broader range of inputs that can shape the final artifact without being visible to the assessor. That includes purchased work, extensive unacknowledged assistance, collusion, template reuse, and differential access to support and resources. These risks are documented in the contract cheating literature and associated sector guidance (Newton, 2018; QAA, 2022), as well as in empirical studies linking outsourcing pressures to student experiences and institutional environments (Bretag et al., 2018).

A second, older point is that assessment design powerfully shapes student effort and therefore teaching practice. Constructive alignment work emphasizes that what is assessed drives what students do (Biggs, 1996), and synthesis work on assessment conditions similarly argues that assessment structures shape study behavior in predictable ways (Gibbs & Simpson, 2004). If the dominant tasks reward the production of polished artifacts with limited process accountability, students rationally optimize for product quality rather than durable capability.

Generative AI changes the scale and the baseline. It reduces the time, effort, and skill needed to produce plausible outputs, and it does so for a much larger share of students. When a technology makes it cheap to manufacture the graded artifact, it becomes weaker evidence unless the task is redesigned to preserve a defensible link between the work submitted and the student capability being inferred (Kane, 2013; Messick, 1995). Contemporary work on authentic assessment reinforces the point that defensible evidence depends on task properties, including fidelity to real performance demands and visibility into what the learner actually did (Gulikers et al., 2007; Vlachopoulos & Makri, 2024).

A useful way to express this, without moralizing, is: if AI can do the assessment, it likely was not a good assessment anyway**.** The claim is not that students are "bad." It is that the interpretation and use argument for the score becomes harder to defend when AI expands the set of plausible

alternative explanations for performance. In that framing, the institutional challenge is primarily a validity and evidence problem, with integrity as one important component. A predictable response is to retreat to supervised written exams, but that mostly changes where the AI influence shows up, and it often narrows what the assessment can credibly claim to measure.

1.1.1 The invigilated exam fallback does not restore evidence

A common institutional response to generative AI is to treat supervised, in-person written exams as an integrity fix. This response is understandable because invigilation strengthens identity assurance and blocks real-time tool use. However, it does not make assessment "AI-proof," and it does not resolve the deeper validity problem. In an AI-rich environment, invigilation mainly shifts tool influence upstream into preparation and rehearsal, while narrowing what assessment can credibly claim about higher-order learning (TEQSA, 2025).

The core issue is displacement. Students can use generative systems before the exam to produce condensed notes, model solutions, canonical phrasing, and stepwise procedures aligned to predictable question types. They can then memorize and rehearse these outputs. The assessment environment may authenticate who wrote on the day, but it cannot reliably distinguish durable understanding from rehearsed reproduction of AI-mediated templates. From a validity perspective, the question is not simply whether AI was excluded during the sitting. The question is whether the observed performance supports the intended interpretive argument about what the student knows and can do (Kane, 2013).

This also revives a familiar critique of exam-dominant regimes. Many timed written exams reward recall, speed, and compliance with stable formats more than conceptual transfer and judgment. Generative AI intensifies that critique because it excels at fluent retrieval and template completion. To be clear, some exams can be redesigned to test reasoning through novel prompts, justification, and trade-offs. But many exam systems at scale rely on predictable formats, past-paper cultures, and marking practices that reward recognizable structure, which increases the value of rehearsal and reduces the evidentiary value of the final script.

Within FARABI terms, invigilated exams may reduce risk on some "conditions" dimensions but do not automatically reduce likelihood overall. AI capability fit can remain high in preparation, task predictability is often substantial, and process observability beyond the sitting remains limited. More broadly, a switch in format does not, by itself, validate the interpretation and use of scores. Validity depends on the full inference chain, including construct representation, consequences, and fairness (AERA et al., 2014). The implication is not that in-person assessment is obsolete, but that it is not a universal remedy. Defensible assessment in the AI era depends less on excluding tools and more on designing tasks and portfolios that make reasoning visible and that can evaluate accountable tool use where it is relevant.

1.2 FARABI as integrity triage and curricular diagnostic

If the problem is system-level credibility, institutions need a way to prioritize, not just react. FARABI is positioned as a practical, risk-based framework for assessing the integrity exposure of an assessment portfolio and prioritizing redesign where it matters most. Its first function is stabilization: making risk legible and comparable so that controls are proportionate rather than driven by panic. Its second function is diagnostic. The assessments that remain relatively safer under AI are not randomly distributed. They cluster around properties that imply a different educational model: applied judgment, observable process, collaboration under accountability, and transparent tool use. In that sense, FARABI is a 'Trojan horse'. It is an integrity framework that also exposes curricular misalignment between stated learning outcomes and the evidence actually collected to support them.

This paper is conceptual, but it is not purely speculative. We have used the thinking presented in this paper and applied the FARABI logic to build an AI-supported decision workflow that guides program teams through portfolio triage, generates structured prompts for scoring and justification, and compiles a transparent record of decisions and control choices. The intent is not automation of academic judgment, but reduction of transaction costs in redesign work and improvement of consistency and documentation across programs.

1.3 Why the sector struggles to act decisively, even when it "knows the problem"

Most universities now recognize that generative AI has changed the assessment environment. What is harder is acting at the pace and scale the situation demands. This is not mainly a knowledge problem. It is a response capacity problem shaped by workload, incentives, and the psychology of slow systemic threats. There are three aspects to this.

Meaningful assessment reform is design-heavy work. It requires coordination across courses, agreement on what evidence is sufficient for key learning outcomes, and changes that affect workload, marking practices, and student experience. These are costly changes, and they rarely fit neatly into existing cycles of validation, accreditation, and quality assurance. As a result, institutions often default to responses that are easier to issue centrally, such as policy statements, guidance notes, and occasional enforcement actions, even when these do not change the evidentiary basis of assessment.

Secondly, institutions operate in accountability environments that reward auditable activity. Compliance legibility can become a substitute for learning evidence. A policy document, an integrity module, or a detection procurement decision can be shown to boards and regulators quickly. Assessment redesign, by contrast, is slower, harder to summarize, and less easily "proven" through paperwork.

Finally, psychology amplifies delay. Systemic threats that are diffuse and cumulative often fail to trigger proportionate attention. Slovic's account of "psychic numbing" and Kahneman's evidence on scope insensitivity help explain why large-scale risks can feel abstract even when they are acknowledged and visible (Slovic, 2007; Kahneman & Knetsch, 1992). Section 6 returns to these

mechanisms and to the inaction and diffidence they produce, because the pattern is structurally similar.

This is the action gap FARABI targets. It is designed for real institutions with limited capacity and mixed incentives. FARABI does not assume perfect compliance, review or universal redesign. It provides a triage logic for deciding where assessment evidence is most exposed, where the stakes are highest, and where minimal, proportionate controls can restore a defensible interpretation and use argument for grades and credentials.

This paper also recognizes a practical constraint: even when institutions agree that generative AI has changed assessment, they often struggle to respond with portfolio-level redesign. The reasons are not limited to policy uncertainty or technical capability. They include workload and coordination costs, incentives that favor auditable activity over evidentiary strength, and predictable patterns of under-response to diffuse systemic threats noted here. Section 6 returns to these dynamics in detail, because without an explanation of the action gap, a risk-based framework like FARABI can easily be misunderstood as a purely technical fix rather than a tool designed for real institutional conditions.

2. What changed: capability shock, scale effects, normalization, and a fast moving policy and market landscape (2024–2025)

The assessment problem did not begin with generative AI, but the conditions that made weak assessments survivable have changed. What shifted is the combination of (a) rapid capability gains in broadly accessible tools, (b) scale effects that turn isolated misuse into population level noise, (c) normalization of routine AI use by students and staff, and (d) a policy and market landscape that is pushing institutions to "adopt and govern" rather than "ban and revert." In practice, this means higher education is now operating inside a new equilibrium where students can cheaply produce plausible assessed artifacts, institutions cannot reliably infer learning from the product alone, and major technology providers are actively building AI into the everyday infrastructure of learning.

2.1 Capability shock plus scale effects: why "a few bad actors" is no longer an adequate mental model

Generative AI is often framed as a cheating problem, but its more consequential feature is low-friction, high-quality production at scale. The economics matter: when producing a polished essay, report, coding solution, presentation narrative, or reflective commentary becomes near-zero cost, the traditional safeguards of effort, time, and skill cease to function as implicit authentication. The result is not simply more misconduct. It is a collapse in the evidentiary value of product-only coursework as a basis for interpreting student achievement.

Evidence from sector surveys supports this shift from exceptional to routine. In the UK, HEPI reports that student use of generative AI is now close to universal, and that use for assessed work is widespread (HEPI, 2025). The DEC's global student survey similarly indicates that the

majority of students are already using AI for study tasks and are doing so in conditions of mixed clarity about what is permitted (DEC, 2024).

This is the first reason the credibility problem for higher education is system-level. Even if only a minority were intentionally outsourcing whole assessments, the mere availability of high-quality automated production changes the meaning of the observed artifact. In validity terms, the question becomes whether the institution can still sustain a defensible interpretation and use argument from the evidence it collects. If the same artifact can be produced by a student with deep understanding or by a student with shallow understanding plus tool fluency, the inference weakens unless the assessment is redesigned to preserve discriminating evidence in favor of demonstrating deep understanding.

2.2 Normalization and student uncertainty: usage is high, but so are integrity and epistemic anxieties

The next shift is that AI is not only being used covertly. It is being used as a routine study aid, and students are asking institutions to clarify the rules and provide competence building. DEC's student survey reports substantial student appetite for clearer guidance and support, including training on responsible use (DEC, 2024). The same survey also shows that students' concerns are not limited to misconduct. They include worries about biased outputs, dependence, and unfairness created by uneven access or uneven knowledge of "how to use the tools well" (DEC, 2024).

This matters for FARABI because it demonstrates that the integrity challenge is paired with a validity challenge. Students are using AI in ways that can blur authorship and inflate performance, but many also doubt their peers' behavior and their own ability to evaluate AI outputs. That combination increases the risk that assessment evidence becomes simultaneously less authentic and less trustworthy.

OUP's "AI Native Generation" work reinforces this point from a slightly different angle: young people report heavy use, but also report uncertainty about truth, bias, and trustworthiness online, and a desire for support in judging reliability (Oxford University Press, 2025).

This is not a marginal issue. It suggests that in an AI-saturated information environment, institutions that ignore AI literacy are not simply missing a contemporary skill. They are weakening the conditions required for reliable academic judgment for all parties.

2.3 Staff readiness and institutional capacity: adoption is growing, but confidence and governance lag

Faculty are not operating from a position of clarity either. DEC's global faculty survey reports substantial faculty use of AI tools, coupled with concerns about students' ability to critically evaluate AI-generated outputs and continued uncertainty about institutional guidance (DEC, 2025a). Put bluntly, this is the capability gap that makes "policy-only" responses fail. If frontline educators are improvising local norms while students are improvising local practice, the

institution is not governing a stable system. It is accumulating inconsistent assessment conditions

At the same time, the response space is constrained by workload and systems. Portfolio redesign, richer process evidence, more observation, and more dialogic assessment all require time, training, and assessment transparency across programs. Institutions that are already stretched tend to default to the least costly response, which is policing. The problem is that the artifact is the wrong control point when the inference is itself unstable.

This is where the validity framing of FARABI becomes practically useful: it shifts the institutional question from "How do we catch misuse?" to "What evidence do we need to support the claims we make about learning, and are we collecting it?"

2.4 Policy landscape: a drift from "ban and detect" toward "redesign and build literacy," with integrity positioned as governance

The 2024–2025 policy landscape is best characterized as convergence without uniformity. Different jurisdictions and agencies vary in emphasis, but many are converging on three ideas: redesign assessment, strengthen AI literacy, and treat governance as an institutional responsibility rather than an individual moral failing.

At the EU level, the AI Act explicitly introduces expectations around AI literacy, including that providers and deployers take measures to ensure a sufficient level of AI literacy for relevant actors (European Union, 2024). The European Commission and OECD's draft AI literacy framework for school education is not targeted at higher education but its logic is relevant. It frames AI literacy as a set of competences across domains such as engaging with AI, creating with AI, managing AI, and designing AI, and it links literacy to ethical and societal understanding (European Commission & OECD, 2025). This framing is aligned with FARABI's emphasis on conditions and observability. When AI is a normal tool, integrity depends on accountable use and on learners being able to interrogate outputs, not merely on prohibition.

OECD work adds another layer by positioning generative AI as part of broader system pressures shaping education, with implications for skills, curriculum, and assessment. The OECD's "Trends Shaping Education 2025" report explicitly treats AI as a major force affecting education systems, reinforcing that institutions should plan for structural change rather than episodic disruption (OECD, 2025a). OECD's AI capability measurement work also signals a shift toward describing AI in terms of functional capabilities and tasks, which is precisely what FARABI operationalizes at the level of assessment design (OECD, 2025b).

Regulatory guidance sits inside this wider convergence, but it should not be treated as the only anchor. The more important point for this paper is that the "center of gravity" in thoughtful guidance has rapidly moved away from detection as the primary solution and toward redesign and literacy as the basis for sustainable integrity.

2.5 The market reality: platform vendors are embedding AI into the learning stack

A final change is that higher education is not responding to AI in a vacuum. It is responding while major technology providers compete to become the default infrastructure for AI-mediated learning and work.

OpenAI's launch of ChatGPT Edu and subsequent large-scale deployments show how quickly institutions can move from experimentation to enterprise rollout (OpenAI, 2024, 2025a, 2025b). Google is positioning Gemini as part of its education ecosystem, alongside training and product integration through Google for Education (Google, 2025). Microsoft is pursuing a parallel path through Copilot offerings and education-focused workflows embedded in Microsoft 365 and Teams (Microsoft, 2024). Meta's "Meta for Education" illustrates that this competition is not only about text generation. It is also about immersive platforms and managed ecosystems marketed directly to educational institutions (Meta, 2025).

This matters for assessment because it makes "going back" implausible. Once AI tools are bundled into institutional logins, productivity suites, learning platforms, and everyday workflows, they become ambient. That pushes institutions toward a choice: either treat AI as a hidden contaminant to be policed, which creates an escalating compliance burden with limited credibility, or treat AI as a normal tool to be governed, which forces assessment redesign and explicit standards for accountable use.

2.6 Why these shifts point to FARABI: the problem is now evidentiary, not only behavioral

Taken together, capability gains, scale effects, normalization, policy convergence, and platform embedding produce the core condition this paper addresses: traditional coursework often no longer yields defensible evidence of higher-order learning without redesign**.** The question is not whether students should or should not use AI. The question is whether institutions can still justify the claims they make about graduates based on the evidence their assessments generate.

This is exactly the gap FARABI is designed to address. FARABI does not assume one universal fix. It treats assessment integrity as a risk triage problem across an assessment portfolio, using a small set of likelihood and impact factors to identify where the inference is most threatened and where redesign yields the largest integrity and validity returns. The next section therefore turns from context to diagnosis: assessment as an evidence problem, framed through validity, interpretation, and use.

3. The assessment problem is an evidence problem

Generative AI forces a basic question that higher education can no longer treat as a niche integrity issue: what exactly does a piece of assessed work count as evidence of? If an assessment score is used to make claims about learning, those claims need a defensible interpretation and use argument, meaning a clear chain from task performance to the intended construct, and from that construct to the decisions made with the result (AERA et al., 2014; Kane, 2013; Messick, 1995).

That framing matters because the generative AI challenge is not mainly about catching individual cases. It is about system conditions. When a tool becomes widely accessible and capable enough to produce plausible course outputs on demand, the probability that a submitted product reflects a student's own reasoning declines unless the task design produces stronger evidence. In validity terms, the threat is not simply "cheating." It is weakened inference.

3.1 Why generative AI changes the inference problem (even when "misconduct" is not proved)

In the traditional coursework model, many assessments rely on take-home artifacts graded mainly on the final product, with limited visibility into how the product was produced. The design assumes that the product is a reasonable proxy for the student's independent competence.

Generative AI changes those assumptions fundamentally. The core change is economic and probabilistic: it becomes cheap to manufacture an acceptable product, quickly, across many common genres of university assessment. Under those conditions, a large share of submissions become epistemically ambiguous from the institution's point of view. That is true even if the institution cannot, and should not try to, prove misconduct in individual cases.

Student and faculty evidence points in the same direction: high expected use, mixed confidence, and demand for clearer rules.

- In the DEC student survey, 59% of students expected their university to increase AI use in teaching and learning, while students also reported concerns that map directly onto validity and legitimacy issues: privacy and data security (61%), trustworthiness of AI-generated content (51%), and bias and fairness in AI evaluations and decisions (32%) (DEC, 2024).
- On fairness specifically, 60% reported worry about fairness of AI in evaluating student work, and the distribution was heavily skewed toward concern (DEC, 2024).
- In the DEC faculty survey (1,681 respondents across 28 countries), about half of faculty agreed they would need to redesign current assignments to make them more AI resistant, while only 5% perceived fully comprehensive institutional AI guidelines in teaching (DEC, 2025a).

This is exactly what a credibility problem looks like. When the system's dominant evidence preference is a take-home product that can be cheaply generated, and when rules and shared expectations remain underdeveloped, the institution's confidence in its own inferences becomes fragile. That fragility does not get fixed by more "integrity messaging." It gets fixed by assessment designs that restore a defensible link between performance and the intended learning claims.

3.2 The pre-AI baseline was already noisy, and AI accelerates the noise

Pre-AI coursework was never a golden age. Product-only assessment has long been vulnerable to influences such as peer over-assistance, ghostwriting, and contract cheating, as well as unequal access to coaching, time, and social capital. Reform arguments existed well before generative AI, including calls to shift assessment toward designs that better support learning while providing stronger assurance that outcomes have been achieved.

Generative AI accelerates the problem because it expands both the accessibility and the quality of non-authored assistance, and it does so across many disciplines and task genres. The key point is not that "students will cheat more." The key point is that the evidentiary value of the final product falls unless the design captures that which the institution claims to be measuring. Contemporary commentary in the field increasingly treats this as a systemic assessment issue rather than a discrete misconduct issue (Molina & Medina, 2024).

The core validity claim of this paper is straightforward. If AI can complete an assessment to the required standard, the assessment is often weak evidence for the intended higher-order learning claim. This is not a blanket dismissal of foundational practice tasks. It is a claim about the evidentiary limits of product-only coursework under conditions of cheap generation when results are used for high-stakes judgments about individual competence. Under widespread AI access, many traditional coursework formats drift toward measuring the capacity to submit a plausible artifact rather than demonstrating intended learning, unless assessments are redesigned so that judgment, reasoning, and accountable tool use are observable and evaluable.

3.3 What "stronger evidence" looks like under AI conditions

Policy work on assessment reform is converging on a simple idea: assessment must continue to serve a dual purpose, supporting meaningful learning while also providing assurance that learning outcomes have been achieved (TEQSA, 2025). Under generative AI conditions, "stronger evidence" usually comes from shifting at least part of the grading basis away from a single polished product and toward evidence that is harder to substitute with generic generation.

3.4 Design responses aligned with AI literacy frameworks

Contemporary guidance converges on a small set of assessment design responses that address AI-mediated validity threats while remaining consistent with AI literacy frameworks. Across contexts, four approaches recur because they strengthen the evidentiary basis for interpreting student performance as evidence of the intended construct, rather than as evidence of successful artifact production under low-visibility conditions. These are discussed below.

3.4.1. Make judgment visible rather than privileging output

A first design response is to require students to surface and justify the judgments that sit behind the final product. This can include defensible choices about assumptions, methods, tradeoffs, and interpretations, alongside a rationale for why these choices fit the task context. The validity

benefit is straightforward: the assessment shifts from evaluating the plausibility of an output to evaluating the quality of reasoning and decision-making that produced it.

3.4.2. Increase observability of process through structured evidence

A second response is to increase the observability of the development process by embedding checkpoints that yield traceable evidence of work-in-progress. Common mechanisms include staged drafts, decision logs, in-class working periods, short oral defenses, supervised problem-framing, and structured reflection tied to specific revisions. The purpose is not surveillance as an end in itself. It is to reduce reliance on a single final artifact as the primary basis for inference, thereby strengthening the chain of evidence supporting claims about individual competence.

3.4.3. Raise task specificity and transfer demands

A third response is to increase task specificity and demand for transfer by anchoring assessments in local data, lived cases, or course-specific materials. This does not prevent AI assistance, but it increases the evidentiary threshold: students must demonstrate situated judgment and contextual application rather than generic completion. In validity terms, the task becomes more discriminating with respect to the targeted construct, because success depends on interpretation and adaptation, not merely fluent production.

3.4.4. Make tool use accountable through disclosure and verification expectations

A fourth response is to make tool use accountable where AI is permitted. This typically involves requiring students to disclose how tools were used, what outputs were verified, and what steps were taken to evaluate accuracy, bias, and contextual fit. This aligns with emerging AI literacy approaches that emphasize critical evaluation and responsible use, rather than a simple binary of permission versus prohibition (European Commission, 2025; Oxford University Press, 2025).

3.4.5 Governance considerations and an equity caution

A governance caution follows directly from the validity logic. Interventions that appear compelling in principle, particularly detection-first approaches or highly visible “security” measures, can impose equity costs and still fail to restore validity if they do not improve the quality of evidence supporting the intended learning claims. The aim is not to maximize control. The aim is to protect the assessment interpretation and use argument by ensuring that the evidentiary basis for high-stakes decisions remains credible under AI conditions.

3.5 Mapping validity threats to FARABI’s likelihood subfactors

FARABI’s likelihood dimension operationalizes the validity problem outlined above by translating general threats to inference into four portfolio-comparable questions. Rather than treating “AI risk” as a diffuse concern, it specifies the conditions under which an assessment’s evidentiary value is most likely to degrade when generative tools are readily available.

3.5.1. AI capability fit

The first question concerns the degree of alignment between readily available AI tools and the task's core demands, including genre requirements, depth of reasoning, calculation, coding, and synthesis. As AI capabilities expand and access frictions decline, a growing proportion of conventional coursework tasks move into a "high fit" category unless they require forms of situated judgment that cannot be readily delegated to a tool (OECD, 2025a, 2025b).

3.5.2. Task predictability

The second question concerns the predictability of the task: the extent to which prompts are templated and outputs are stable. Highly predictable tasks lower the marginal cost of completion via prompt engineering, recycling, and automated drafting. Predictability does not, by itself, imply low educational quality. However, where process evidence is limited, it reduces the evidentiary value of the final product for making individual-level inferences about competence.

3.5.3. Process observability

The third question concerns observability: how much of the student's work is visible to the assessor. When instructors see only the final artifact, interpretation depends heavily on assumptions about authorship and independent production. When structured process evidence is available, the assessment can accommodate tool assistance while still supporting a defensible claim about the student's reasoning, decision-making, and learning.

3.5.4. Assessment conditions

The fourth question concerns the conditions under which the work is produced and submitted, including identity assurance, time windows, collaboration rules, tool rules, and supervision arrangements. These conditions shape the ease with which a product can be manufactured and, correspondingly, the confidence with which the institution can interpret performance as evidence of the intended construct.

3.5.5 Bridging validity theory and triage

Taken together, these four questions bridge validity theory and FARABI as a practical triage framework. The central issue is not whether AI is present, but whether the assessment evidence still warrants the interpretations institutions attach to scores and credentials under contemporary technological conditions. Section 4 turns this validity argument into an operational model by defining the FARABI dimensions and their use in portfolio-level redesign decisions.

4. FARABI: A risk-based framework for AI-resilient assessment portfolios

Section 3 framed generative AI as a validity shock: when assessment relies heavily on take-home products with limited visibility into how work is produced, the evidentiary value of the final

artifact weakens under conditions of cheap, scalable generation. This section introduces FARABI (Framework for AI-Resilient Assessment and Balanced Integrity) as a portfolio-level response. The central premise is that institutions do not need a single universal solution. They need a way to identify where assessment evidence is most exposed, where the consequences of failure are greatest, and which redesign moves most directly strengthen the interpretation and use argument that underwrites academic judgment (AERA et al., 2014; Kane, 2013; Messick, 1995).

FARABI is therefore not presented as a detector regime or a moral stance on student behavior. It is a structured way to translate an abstract validity problem into program-level redesign priorities. This positioning aligns with the direction of contemporary sector guidance, which increasingly treats assessment reform and AI literacy as coupled responsibilities of governance rather than as separable questions of misconduct control (DEC, 2025b; TEQSA, 2025; HEPI, 2025).

4.1 Conceptual scope and boundary conditions

FARABI is designed for program-level use. Its unit of analysis is the assessment portfolio across a degree, not isolated tasks in single units. This choice follows directly from the diagnosis in Section 3. Under AI conditions, vulnerability is uneven. Some tasks are readily substituted by widely available tools, while others remain robust because they make reasoning, contextual judgment, or development visible. A portfolio lens enables programs to concentrate redesign effort where it most improves inferential defensibility and to avoid overcorrecting by hardening every task indiscriminately.

Two boundary conditions follow. First, FARABI is not a claim that invigilated formats are the default remedy. Controlled conditions can improve identity assurance and increase observability at the point of performance, but they do not by themselves secure the intended interpretation and use of scores. As discussed in Section 1.1.1, AI influence can be displaced into preparation and rehearsal, and predictable formats can remain vulnerable even under supervision. Over-reliance on invigilated formats can also narrow construct coverage and shift burden onto students with constrained time and mobility. The point is not to prefer one format. It is to restore evidence while preserving the breadth of learning claims.

Second, FARABI does not treat detection as the primary integrity mechanism. Detection tools may provide a weak signal and a prompt for conversation, but they are not a stable foundation for validity or adjudication at scale. The framework begins instead with redesign and transparency because those are the levers that most directly strengthen the link between performance evidence and intended constructs.

4.2 Risk as a portfolio decision problem

FARABI operationalizes assessment risk as the interaction of two dimensions:

$$\textbf{Risk} = \textbf{Likelihood} \times \textbf{Impact}$$

FARABI is intended as a heuristic decision aid for portfolio triage, not a psychometric measurement instrument. The scores are designed to support comparative prioritization and

structured discussion under uncertainty, not to generate precise estimates of misconduct or learning. Empirically, an important research task is to test inter-rater consistency and whether FARABI-guided redesign improves evidentiary strength and reduces ambiguity in practice (see Section 7).

4.2.1 Likelihood

This captures how exposed a task is as evidence under AI conditions. FARABI specifies four subfactors that, taken together, explain much of the substitution risk for conventional coursework:

- **AI capability fit:** the extent to which widely available tools match the task's core demands (genre, reasoning depth, calculation, coding, synthesis). As AI capability broadens and becomes easier to access, more conventional coursework tasks drift into "high fit" territory unless they require judgment that is not readily outsourced (OECD, 2025b).
- **Task predictability:** the degree to which the task is templated. Predictable prompts with stable output genres invite prompt reuse, prompt engineering, and automated drafting.
- **Process observability:** how much of the student's work is actually seen. When instructors see only a final artifact, inference depends heavily on trust in authorship. When process evidence exists, assessment can tolerate tool assistance while still supporting defensible claims.
- **Assessment conditions:** the surrounding conditions of production and submission (time windows, supervision, identity assurance, collaboration rules, tool rules). Conditions shape how easily a product can be manufactured and how confidently an institution can interpret it. Stronger conditions can reduce some risks by increasing control and observability at the point of performance. However, they do not automatically reduce AI capability fit or task predictability, and a highly predictable exam can remain strongly amenable to AI-supported preparation and rehearsal even under tight supervision, which limits what the resulting performance can evidence with validity.

4.2.2 Impact

Impact captures the consequence of evidence failure. FARABI specifies four subfactors that reflect the decision stakes attached to an assessment:

- Weight in grade and progression: how much the task shapes outcomes for students.
- Centrality to core learning outcomes: whether the task is a primary evidence point for the outcomes the program claims to deliver.
- Transfer risk: the risk of passing students who cannot carry competence into later study, professional practice, or licensure contexts.
- External exposure: the extent to which weak evidence threatens accreditation expectations, professional thresholds, or public trust commitments.

Two implications follow. First, low-likelihood tasks can still be high-impact and therefore require careful design. Second, high-likelihood tasks can be low-impact and may be appropriate places to permit AI-assisted learning without undermining overall program credibility, provided the portfolio contains sufficient high-quality evidence points elsewhere. FARABI therefore encourages broad banding (for example, low to critical) to support prioritization while avoiding false precision.

4.3 Declared AI-use modes and the construct problem

A persistent difficulty for institutions is not merely student behavior, but governance ambiguity. Formal policies often coexist with informal practices negotiated across courses, and the result is inconsistent expectations and unclear constructs. FARABI therefore requires an explicit declaration of the intended role of AI in each assessment. The point is not to prescribe a single institutional policy stance. It is to ensure the assessment design and marking criteria match what is being claimed.

FARABI specifies four modes:

- Human only: AI tools are not used in producing assessed work. This mode is defensible only when the construct requires it and the institution can provide fair conditions for compliance.
- Assist: AI is permitted for bounded support (brainstorming, language editing, translation support, low-stakes feedback), while the student remains responsible for substantive decisions and verification.
- AI–human interface: AI use is integrated into the task and evaluated through disclosure, verification, and critical evaluation. In this mode, the construct shifts toward responsible judgment with tools rather than raw production.
- Autopilot: AI generates substantial portions of the artifact. This mode is defensible only where the learning claim is explicitly about selection, evaluation, and accountability rather than original production.

The theoretical point is straightforward. AI-use policies are not add-ons. They alter what the task evidences. Declaring an AI mode makes the construct explicit and reduces the risk that the institution grades one thing while claiming another.

4.4 Minimal controls as validity supports

FARABI assumes there is no scalable "perfect control." The objective is proportionate: strengthen the evidentiary basis of assessment with the least burdensome set of supports consistent with defensible inference. FARABI therefore groups controls into four families, each aligned to a different weakness in the interpretation and use argument.

Design controls increase construct relevance by requiring contextual judgment. They typically shift tasks away from generic completion and toward situated decision-making, including explicit tradeoff justification, locally anchored data or cases, and staged tasks in which later outputs depend on earlier choices and feedback. These moves do not eliminate AI assistance, but they raise the evidentiary bar from fluency to judgment.

Process controls increase observability. They provide credible visibility into development without requiring continuous surveillance. Examples include structured checkpoints tied to key claims or decisions, brief oral explanations of reasoning, or in-class problem framing and critique sessions that generate traceable work products. The purpose is evidentiary: to reduce dependence on a single polished submission as the sole basis for inference.

Assurance controls strengthen interpretation and adjudication. They include clear disclosure norms, marking criteria that reward verification and judgment rather than surface polish, and defined pathways for addressing ambiguous cases with due process safeguards. Their role is to prevent institutions from treating weak technical signals as decisive proof.

Governance controls operate at portfolio level. They include explicit mapping of assessments to program outcomes to ensure there are sufficient high-observability evidence points across the degree, periodic review triggered by capability changes, and staff development oriented around reusable redesign patterns. This family matters because the dominant institutional failure mode is decoupling: policy evolves faster than practice, and visible compliance substitutes for improved evidence.

4.5 Balanced integrity as an implementation stance

The framework's purpose is easily undermined by predictable implementation failures. One is symbolic adoption, in which FARABI becomes another reporting exercise with little effect on assessment designs. The other is punitive hardening, in which risk is treated as suspicion and controls expand until learning quality and trust are damaged. FARABI therefore incorporates a balanced integrity stance, understood as governance principles that protect validity while constraining compliance drift. Five principles follow as illustrated in Figure 1 and discussed further below.

Figure 1. FARABI Model Principles

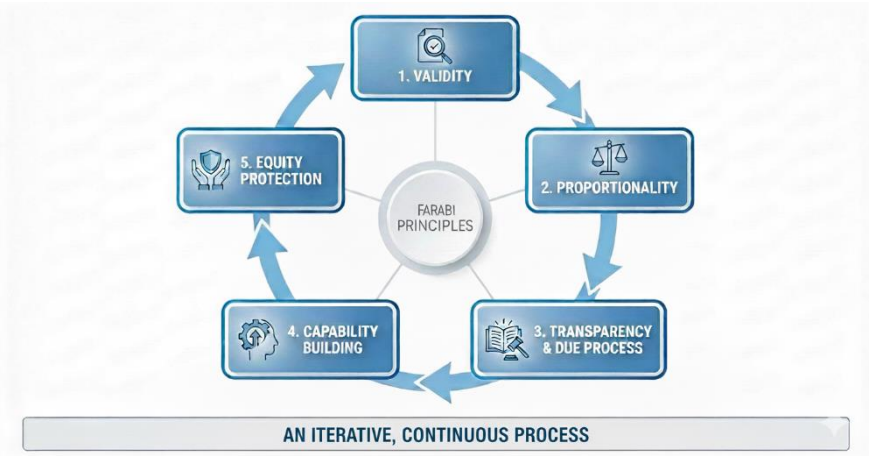

Source. Developed by the author

- Validity primacy: the central question is whether the assessment supports defensible learning claims under current conditions, not whether the institution can prove individual AI use.
- Proportionality: controls should match risk and consequence. Programs should allocate redesign effort where it changes inference most.
- Transparency and due process: ambiguous cases require fair procedures and standards of evidence that are appropriate to the limits of attribution.
- Capability building: credible redesign requires staff time, shared patterns, and institutional support, not only policy statements.
- Equity protection: assessment redesign and integrity controls have distributional consequences; equity is therefore part of validity, not a separate ethical add-on.

4.6 FARABI's intended use and its link to the "Trojan horse" claim

Used as intended, FARABI enables programs to make explicit the evidentiary logic of their assessment portfolios: where learning outcomes are claimed, where evidence is strongest, where AI conditions create vulnerability, and which redesign moves are warranted first. The point is to restore credible inference without collapsing into either detection dependence or across-the-board restriction.

This portfolio logic also sets up the paper's subsequent claim. Once programs consistently prioritize observability, contextual judgment, and accountable tool use to preserve defensible evidence, "AI-resilient assessment" functions as a curriculum filter. It changes what is taught, what is rewarded, and what graduates are expected to demonstrate. Section 5 examines those implications and argues that FARABI therefore operates as a Trojan horse: it enters as an

assessment credibility framework and ends by pushing institutions toward a different educational model.

4.7 Operationalizing FARABI as an AI-supported workflow

A practical advantage of a structured risk framework is that it can be instantiated as a decision-support workflow. We have operationalized FARABI as an AI-assisted process that helps programs (i) elicit the relevant task information, (ii) score likelihood and impact with explicit rationales, (iii) declare an AI-use mode aligned to the intended construct, and (iv) select proportionate control patterns that strengthen evidentiary defensibility. The workflow outputs a program-level portfolio view rather than a course-by-course compliance checklist.

The purpose of this tooling is not to replace academic judgment or to produce mechanically "correct" scores. Its purpose is to standardize the questions asked, surface assumptions, reduce avoidable variation in scoring discussions, and create an auditable rationale for redesign choices. Used in this way, AI becomes part of the response capacity: it supports the design and governance work that institutions struggle to scale, while keeping responsibility for decisions with program leaders and academic governance.

5. The Trojan horse: AI-resilient assessment implies a different educational model

FARABI is presented in Section 4 as a risk-based way to restore defensible assessment evidence. Here we make a second claim, and it is why the framework matters beyond compliance. Once institutions treat assessment as an evidence system under AI conditions, they discover that "safer assessment" is not a technical patch. It quietly forces a shift in what higher education rewards, what it teaches, and what it asks students to show. That is the Trojan horse. Institutions can adopt FARABI to protect credibility, and use of the system can lead to redesigning the curriculum around applied judgment, process evidence, accountable tool use, and discipline-embedded AI literacy.

5.1 Assessment is not neutral: it is a curriculum filter

Assessment does more than measure learning. At scale, it selects what students prioritize, what staff teach explicitly, and what counts as success. Under the traditional coursework model, the dominant signal has often been the polished final artifact. Under generative AI conditions, that signal becomes cheaper to manufacture and therefore less reliable as evidence of the underlying construct.

The institution faces a choice. It can keep grading products while hoping that authorship holds. Or it can redesign tasks so that the graded evidence depends on things that are harder to substitute with generic generation: contextual judgment, justified decisions, verification, and visible development. That second option is not merely an integrity move. It changes what learning looks like. The policy conversation is already drifting in this direction. The DEC's global review of AI-integrated assessment practices frames reform around redesign patterns rather

than enforcement, emphasizing that the "unit of analysis" is the assessment system, not individual misconduct cases. (DEC, 2025b; Molina & Medina, 2025).

5.2 From production to judgment, selection, and accountability

A quiet assumption in many coursework tasks is that competent students can produce the required artifact, and that production itself indicates competence. Generative AI breaks that link for many genres. If a tool can draft, structure, or even solve much of what is being graded, then the educational value of "production" declines as a discriminator, unless the task is redesigned.

What remains scarce is not text. It is judgment. That includes deciding what matters in a situation, selecting and justifying methods, checking claims against evidence, spotting failure modes, and taking responsibility for decisions. This is consistent with a broader argument in the higher education literature that the value proposition shifts toward human capabilities that are not reducible to routine information work, particularly integrative thinking and judgment under uncertainty (Aoun, 2018; Luckin, 2018).

FARABI's AI use modes make this explicit. In "AI–human interface" tasks, the assessed construct is no longer "write the report." It becomes "use tools responsibly to reach a defensible conclusion," including verification, bias awareness, and transparency about tool limits. Section 3 already framed this as a validity move. Here we state the implication. Credible assessment in the AI era pushes universities toward teaching and rewarding accountable judgment, not output fluency.

5.3 Authenticity is not a "vibe": it is an evidentiary property

The term "authentic assessment" is often used loosely, as if realism automatically creates credibility. Under AI conditions, authenticity is not about role-play aesthetics. It is about whether the task generates evidence that supports the institution's learning claims. In practice, that usually means at least one of the following is true:

- The task is anchored in local context, local data, or discipline-specific constraints that require situated judgment.
- The student's reasoning and development are observable enough that the final artifact is not the sole evidence point.
- The task requires defensible tradeoffs and verification steps that cannot be satisfied by generic completion.

This is also where "engagement" narratives can mislead. Institutions are understandably experimenting with AI for advising, tutoring, and engagement, but engagement tools do not automatically improve assessment validity. In some cases, they can accelerate a convenience-first culture where the institution optimizes for smooth production rather than strong evidence. The DEC review of AI for student engagement illustrates how quickly institutions are adopting AI

in engagement functions, which makes it even more important that assessment does not drift into a parallel logic of automation and plausible output. (DEC, 2025b).

5.4 From policing to evidencing learning

FARABI's process and assurance controls imply a different feedback culture. If institutions want credible evidence, they need more structured visibility into development, not more suspicion at the deadline. This does not require turning courses into surveillance. It requires shifting staff effort toward moments that generate informative evidence: problem framing, decision checkpoints, critique, short oral explanations of reasoning, and evaluation of how a student verified and improved work. Under AI conditions, these moments do double duty. They support learning, and they strengthen the interpretation and use argument by making the student's contribution legible. The HEPI report on AI and the future of universities reflects this governance reality. AI affects not just student behavior, but the operating model of institutions, including assessment practices and staff roles. In other words, the "assessment question" is part of a broader institutional redesign problem (HEPI, 2025).

5.5 AI literacy becomes integrity infrastructure, not an elective add-on

If assessment shifts toward accountable tool use and defensible judgment, then AI literacy stops being optional. It becomes part of the integrity infrastructure, because students cannot meet the assessment standard without being able to evaluate AI outputs, identify limitations, and document responsible use.

This aligns with the direction of major literacy initiatives. The joint European Commission and OECD AI literacy work frames AI literacy as a shared responsibility across the education ecosystem, with competence areas that emphasize critical evaluation, responsible use, and understanding limitations (European Commission & OECD, 2025). The Trojan horse claim is therefore concrete: if you redesign assessment so that the evidence is defensible under AI conditions, you will end up teaching AI literacy within disciplines, because it becomes necessary to perform the task, not merely to comply with a policy.

5.6 Platform-scale dynamics make the Trojan horse unavoidable

The assessment conversation often treats AI as a tool students independently adopt. That framing is already obsolete. AI is becoming a platform layer across the education sector, which changes incentives and default behaviors.

- Google's "AI and the Future of Learning" positions AI as part of a long-term learning infrastructure, explicitly linking AI capabilities to learning science and broad deployment (Google, 2025).
- OpenAI has explicitly moved into the institutional market with ChatGPT Edu, designed for university-wide deployment with enterprise controls (OpenAI, 2024).

- Microsoft is bundling “Copilot Chat” and related AI capabilities into education offerings, normalizing AI use inside the productivity suite that already structures much student work (Microsoft, 2024).
- Meta is also positioning itself in education through “Meta for Education” initiatives and related programs, signaling that the platform contest extends well beyond one vendor (Meta, 2025).

The implication is clear. Even if a university wanted to treat AI as an external disturbance, platform integration will make AI a default feature of the environment in which academic work is produced. That is why the credibility problem is system-level, and why a framework like FARABI is better understood as institutional adaptation, not a policy memo.

5.7 What FARABI implies, stated plainly

That is not a small tweak. It is a shift in what higher education claims to cultivate. This is why FARABI is a Trojan horse: it enters as a credibility fix and exits as a curriculum and governance reform program. If the implications are this clear, why is institutional response so often restricted to partial measures and symbolic policies? We argue that the limiting factor is not awareness of the problem but the political economy of reform: constrained capacity, misaligned incentives, and audit logics that privilege procedurally legible controls over redesign that restores the evidentiary basis for interpretation and use. Section 6 specifies these mechanisms.

6. Why higher education struggles to act (and how FARABI avoids becoming a compliance factory)

If the diagnosis above is broadly accepted, then the obvious question is why institutional response remains uneven and often superficial. The simplest answer is not ignorance. Most universities understand the problem. The constraint is that credible reform is design-heavy, politically awkward, and poorly rewarded by existing governance systems. Under those conditions, many institutions default to symbolic responses: statements, workshops, and selective rule tightening, while the underlying assessment portfolio changes slowly.

This section makes a skeptical claim. Higher education is currently better at producing policies than producing new evidence and assessment systems. That gap is not primarily a moral failure. It is a predictable outcome of capacity limits, incentive structures, and human psychology.

6.1 Recognition without readiness

General AI awareness has risen quickly, but operational readiness has not, at least to the same extent. Institutions face three practical constraints.

6.1.1 Staff capability and time.

Redesigning assessment is skilled work. It requires discipline-specific thinking about constructs, task design, and evidence, plus the ability to supervise process without inflating workload

beyond what teaching contracts allow. Surveys and sector reporting consistently suggest many instructors expect to redesign tasks, but do not feel they have clear, stable institutional guidance or sufficient support to do it well (DEC, 2025a; HEPI, 2025).

6.1.2 Infrastructure mismatch

Many institutions still operate with learning technologies built around submission, grading, and plagiarism checking for product-based coursework. Those systems are not designed to capture richer process evidence or structured accountability for tool use. Adding AI into that environment tends to increase ambiguity rather than resolve it.

6.1.3 Governance Fragmentation

Responsibility is diffused across academic integrity units, teaching and learning centers, IT, program committees, and senior management. The result is predictable: high activity, low coherence. Institutions produce guidance, but programs and faculty are left to interpret it task-by-task, often inconsistently.

6.2 Design work matters and is expensive

FARABI pushes toward redesign because that is where validity is restored. Unfortunately, redesign is exactly the kind of work universities are structurally bad at scaling. Portfolio reform requires programs to answer uncomfortable questions:

- Which assessments actually evidence our outcomes?
- Where is process visible?
- Which tasks are high impact and high vulnerability?

Those conversations are time-consuming and politically sensitive because they imply that some long-standing assessment practices are weak evidence.

The temptation is to choose cheaper substitutes: new wording in integrity policies, more warnings, more detection, or tighter restrictions. These moves are legible, fast, and centrally controlled. They are also often weak responses to an evidence problem. Detection is limited as proof, restrictions are hard to enforce at scale, and both can damage trust when they create the sense that students are presumed guilty (Turnitin, 2025; TEQSA, 2025).

6.3 The action gap is psychological, not just organizational

There is a deeper reason universities underreact to system-level assessment risk: humans are not built to respond well to slow, diffuse, high-consequence threats. This is where the climate parallel is not rhetorical. It is structurally informative. Two mechanisms matter.

6.3.1 Psychic numbing.

As harms become large and abstract, people feel less, not more. Risk becomes statistics rather than urgency, which reduces motivation to act (Slovic, 2007). In the assessment context, this shows up as a strange calm. Institutions can acknowledge large-scale credibility erosion while behaving as if it is a series of isolated incidents.

6.3.2 Scope insensitivity.

People respond weakly to scale. A problem affecting 5 students and a problem affecting 5,000 can trigger similar emotional and managerial reactions, especially when the costs of response are immediate and the benefits are uncertain (Kahneman, 2011). Under generative AI, the credible baseline threat is not one student using a tool. It is that product-only evidence becomes systematically ambiguous. But the institutional response often remains calibrated to the old misconduct frame. An analogy with climate change is useful here (Weber, 2006; Gifford, 2011). Universities recognize the direction of travel, but delay large redesign because it is costly now, contested now, and the consequences feel distributed later.

6.4 Incentives, audit logic, and the comfort of legible procedures

Even when leaders want meaningful reform, institutional incentives push toward responses that are easy to document. Universities are governed through auditability. What gets rewarded is what can be shown, such as policies, procedures, training completion, committee minutes, and system settings. These are not worthless. But they are not the same as improved validity.

This is a classic decoupling pattern: formal structures signal control and legitimacy, while day-to-day practice changes more slowly (Meyer & Rowan, 1977; Bromley & Powell, 2012). Under pressure, institutions also imitate each other's visible responses, which can produce fast convergence on policy language without equivalent convergence on assessment redesign (DiMaggio & Powell, 1983). Audit dynamics amplify this by privileging what is countable and reportable over what is educationally meaningful but harder to verify (Power, 1997).

Generative AI intensifies the problem because the "right" response is messy. Strong evidence often requires richer process, contextual judgment, and staff time. Those are harder to audit than a policy update or a new tool procurement.

6.5 Identity, fairness, and the politics of changing the rules

Assessment reform is also a cultural and political challenge inside institutions. Faculty identity is tied to disciplinary norms and autonomy. Redesign can feel like an accusation that prior practice was inadequate. In some departments, it also raises status concerns: moving away from traditional long-form essays or problem sets can be misread as lowering standards, even when the goal is stronger evidence. Students, meanwhile, experience rules about AI as unstable and sometimes inconsistent across courses. Where expectations are unclear, students fill the gap with peer norms and platform defaults, which are rarely aligned to institutional validity goals (DEC, 2024; HEPI, 2025).

Overreliance on surveillance-like controls, blanket restrictions, or punitive inference from weak evidence can produce legitimacy crises of its own. This is why FARABI's emphasis on due process and proportionality is not moral grandstanding. It is a stability requirement for governance.

6.6 How to deploy FARABI without building a compliance factory

A risk framework can easily become busywork. If FARABI is implemented as another form to complete, it will fail. The deployment logic has to match the reality described above. Four principles keep FARABI from turning into compliance theater:

- Use FARABI at program level, not course-by-course paperwork. The purpose is to re-balance the portfolio so that high impact learning claims are backed by sufficient observable evidence points. Programs can tolerate lower-observability tasks when other assessments do the heavy lifting.
- Concentrate redesign where it changes inference. FARABI is a triage tool. High likelihood plus high impact tasks are redesign priorities. Low impact tasks can often remain AI-permissive and learning-oriented, provided the portfolio as a whole remains credible.
- Stop treating detection as the core control. Detection can be one input into a conversation, but it is not a reliable foundation for validity or adjudication at scale (Turnitin, 2025; TEQSA, 2025). FARABI is built around redesign and accountable tool use because those are more durable.
- Build shared patterns, not bespoke heroics. The scalable unit is not a perfect assignment. It is a small set of reusable design patterns: structured checkpoints, short oral reasoning checks, transparent AI-use disclosures, and rubrics that reward verification and judgment. Institutions should invest in pattern libraries and staff development that helps departments adapt patterns to their discipline, rather than insisting on one universal model.

We have been direct in this section for a reason. If universities treat generative AI as an integrity nuisance rather than a validity shock, they will keep producing centrally legible activity while assessment credibility erodes. FARABI is intended to make the real work unavoidable, while still keeping it proportionate. If FARABI is the mechanism, what are the broader implications for institutional practice and research? The discussion below develops the deployment implications, including equity risks, AI literacy as integrity infrastructure, and a research agenda that measures whether reform improves evidence rather than merely increasing compliance.

7. Implications and research agenda

The preceding sections argue that generative AI weakens the evidentiary value of many product-focused coursework tasks, and that the appropriate institutional response is portfolio redesign rather than detection-led enforcement. Section 7 draws out the implications of that claim. The intent is not to offer a procedural blueprint. It is to specify what follows, at the level of

governance, equity, and research, if higher education treats assessment primarily as a validity problem and FARABI as a decision aid for rebuilding defensible evidence.

7.1 Implementation implications: portfolio logic and institutional capacity

A first implication is that reform should be organized at the program level. Assessment credibility is not secured by optimizing individual tasks in isolation. It is secured when a program can point to a coherent portfolio in which high-stakes learning claims are backed by sufficient evidence points with meaningful observability of reasoning, decision-making, and verification. A portfolio lens also reduces the risk of disproportionate controls: programs can tolerate lower-observability tasks in some areas when other assessments carry the evidentiary load for central outcomes.

This has a direct capacity rationale. Redesign is skilled, time-intensive work, and universities cannot plausibly redesign everything at once. A risk-based triage approach allows institutions to concentrate effort where the marginal validity gain is highest: assessments that are simultaneously high impact and highly vulnerable under current AI conditions. In this sense, FARABI functions as a prioritization device that links institutional constraints (time, staff capability, governance coordination) to an explicit validity goal, thereby reducing the likelihood that response effort is absorbed by symbolic policy activity.

A second implementation implication concerns governance coherence. The AI environment makes fragmented policy particularly costly. When institutional guidance is interpreted locally on a course-by-course basis, student expectations become inconsistent and staff practice becomes uneven, which increases both legitimacy risk and dispute likelihood. The sector-level pattern is already visible: rapid adoption and experimentation, accompanied by uncertainty about rules, acceptable assistance, and evidentiary standards for adjudication (DEC, 2024; HEPI, 2025). A framework like FARABI does not eliminate discretion, but it can discipline it by requiring programs to declare constructs and AI-use modes, and to justify controls as validity supports rather than as generalized suspicion.

7.2 Equity as a validity condition: distributional consequences of integrity controls

A second implication is that integrity responses have distributional effects, and those effects are not peripheral to assessment credibility. Validity arguments are weakened when assessment conditions systematically disadvantage certain groups or when compliance with rules is unrealistic for students with constrained time, mobility, caregiving responsibilities, or unequal access to quiet space and stable connectivity. Under AI conditions, the most common institutional reflex, increasing supervision and tightening rules, can therefore create a second-order validity problem: the assessment may become more secure while also becoming less fair and less representative of the intended construct for some students. Three equity tensions are especially salient.

First, a shift toward more controlled assessment formats can increase burdens that fall unevenly across student populations. In-person requirements, synchronous checkpoints, and oral

defenses can be appropriate where the construct requires direct evidence of reasoning, but they also impose participation costs that may not be equitably distributed. The relevant question is not whether control increases security. It is whether the new conditions preserve construct coverage and interpretability across groups.

Second, AI itself introduces an access gradient. As tools become embedded within major platforms and productivity ecosystems, the practical boundary between "permitted" and "available" blurs. Some students will have access to more capable tools, paid tiers, and better hardware, and institutional defaults can inadvertently entrench that advantage. If assessments permit AI assistance without addressing access and transparency, performance may partially reflect resource disparities rather than competence.

Third, procedural justice becomes more important as attribution uncertainty rises. If institutions rely heavily on detection indicators or weak attribution signals to infer misconduct, they increase the risk of false positives and contested decisions, and the social costs of error are likely to be uneven. This is not merely a legal concern. It is a legitimacy concern that can undermine trust in assessment decisions as evidence of learning.

The equity implication is therefore not a general appeal to fairness. It is a validity claim: defensible interpretation and use requires that redesigned conditions be evaluated for distributional effects and that adjudication standards reflect the limits of attribution evidence (Molina & Medina, 2025; Reimers et al., 2026). In resource-constrained systems, this constraint is even sharper. High-touch redesign and heavy supervision models assume staffing and infrastructure that may not exist, which makes triage frameworks more relevant, not less.

7.3 AI literacy as integrity infrastructure

A third implication is that AI literacy becomes a component of the assessment evidence system rather than an optional skills layer. If assessments are redesigned so that students are expected to use AI in bounded ways (Assist) or to demonstrate responsible judgment with tools (AI–human interface), then students must be able to evaluate output quality, verify claims, identify limitations and bias, and document their use transparently. Without that competence, assessment becomes either a prohibition regime that is difficult to enforce or an informal "anything goes" environment in which the institution cannot plausibly interpret what student work signifies.

The stronger claim is that literacy is necessarily discipline-situated. The competencies required to verify AI-supported reasoning in law, nursing, engineering, finance, or history differ in methods, evidentiary standards, and failure modes. Generic admonitions to "use AI responsibly" are unlikely to produce the accountability evidence that FARABI's logic requires. What is needed is discipline-embedded literacy tied to concrete practices: verification routines, disclosure norms, critical evaluation of sources and uncertainty, and awareness of how tool use interacts with professional standards.

This implication aligns with emerging literacy frameworks that emphasize critical understanding, responsible use, and awareness of limitations as core components of AI competence (European Commission & OECD, 2025). It also connects to a wider educational argument: as production becomes easier to outsource, the scarce educational target shifts toward metacognition, context-sensitive reasoning, and values-informed judgment, which are central to human intelligence and to professional competence (Luckin, 2018).

7.4 Platform dynamics and institutional strategy

Section 5 argued that platform-scale integration makes AI ambient rather than optional. The implication for institutions is that assessment governance cannot be reduced to student behavior management. It requires institutional strategy: clarity about tool ecosystems, data governance, disclosure expectations, and equitable access, as well as the capability to maintain coherent standards across courses and programs.

This matters because platform embedding changes default behavior. If AI is integrated into the everyday environments where academic work is produced, then prohibition-first policies become less credible, enforcement costs rise, and the gap between formal statements and practice expands. A validity approach therefore implies institutional investment in governance capacities: not only rules, but shared standards of evidence, supported disclosure practices, and assessment designs that remain interpretable in an AI-saturated environment (HEPI, 2025).

This is also where broader policy landscapes enter indirectly. Even when regulation is not education-specific, regulatory expectations around risk governance, competence, and accountability contribute to the institutional environment in which universities operate. The implication for higher education is that credible governance will increasingly require demonstrable capacity to manage AI use responsibly, rather than reliance on assertions of prohibition.

7.5 A research agenda oriented toward evidence, not policy activity

A central risk identified in Section 6 is decoupling: highly visible institutional activity that does not materially strengthen assessment evidence. The research agenda should therefore focus on whether reforms improve the interpretation and use argument in practice, rather than merely documenting that policies exist.

Five research directions follow from FARABI's logic.

- Validity evidence for redesigned assessments under AI conditions. Which redesign features most improve the link between performance and intended constructs, and how stable are these links across disciplines and levels of study? Research should test whether observability and accountability mechanisms improve interpretability rather than simply changing surface features of tasks.

- Observability mechanisms and workload trade-offs. Which approaches to process evidence provide credible visibility without unsustainable staff workload? Comparative

studies should examine redesign patterns not only for learning outcomes but for feasibility, including marking burden, staff time, and the scalability constraints of large programs.

- Equity effects of integrity controls. What are the distributional impacts of common responses such as increased supervision, oral checks, and more controlled conditions? Which combinations protect both evidentiary strength and access? This includes procedural justice research on dispute processes and the evidentiary status of detection indicators.

- Accountable tool use as an assessable construct. If programs choose Assist or AI–human interface modes, what does competence look like, how is it taught, and how is it assessed? Research should focus on verification practices, critical evaluation, disclosure routines, and the ways these practices develop across a degree.

- Implementation under constraint. In resource-limited contexts, what minimal redesign set produces the greatest validity gain per unit of staff effort? The Global South literature emphasizes that many systems cannot adopt high-touch Western reform templates; research should therefore evaluate triage-based models and reusable patterns that create stronger evidence within real capacity limits (Reimers et al., 2026; Molina & Medina, 2025).

7.6 Purpose beyond employability?

Finally, the paper implies a broader question it does not resolve. If AI changes the labor market value of some forms of graduate output, employability becomes a weaker organizing metric for institutional strategy. Universities may need to justify their value more directly in terms of judgment, civic reasoning, ethical competence, and the capacity to work responsibly with powerful systems. FARABI does not answer that question. It identifies the immediate problem: assessment credibility depends on defensible evidence under current technological conditions. The wider question of purpose and value is a fertile area for further work.

Section 7 has therefore specified the implications of the validity framing and sketched a research agenda oriented toward evidence rather than compliance. Section 8 translates the argument into propositions that can be tested and debated.

8. Eight propositions

This paper makes claims that should be falsifiable in practice, not treated as slogans. The propositions below summarize the core argument in Figure 2 in testable form and link each claim to the evidence and theory discussed in Sections 2–7 in the following discussion below.

Figure 2. FARABI: Eight Propositions

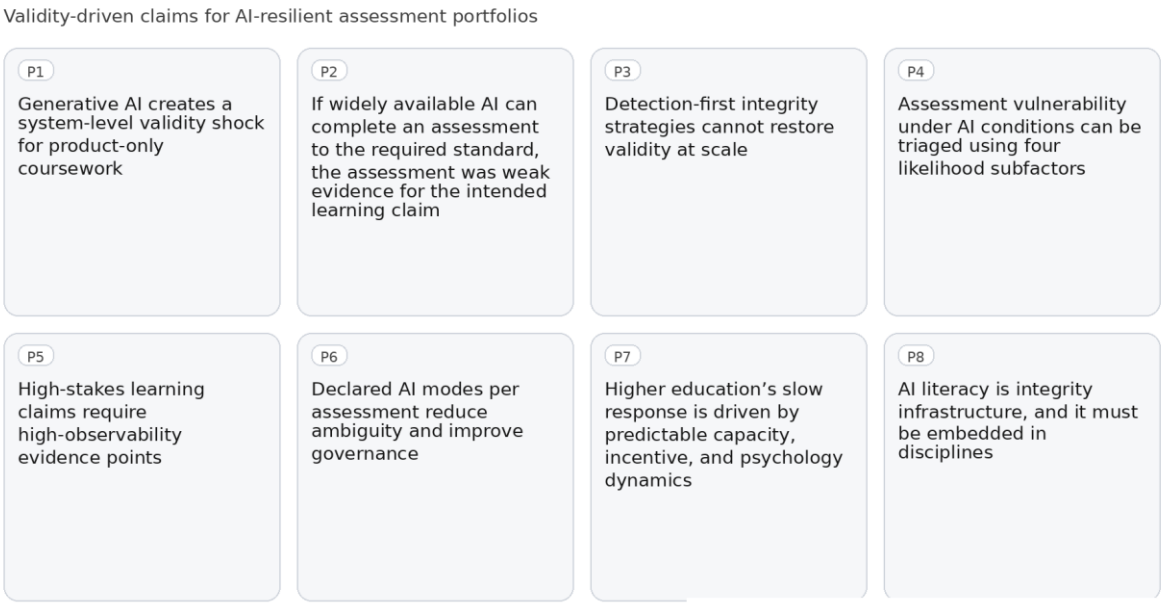

Source. Developed by the author

Proposition 1. Generative AI creates a system-level validity shock for product-only coursework.

When the dominant evidence for learning is a take-home product with low process visibility, widespread access to generative AI reduces the evidentiary value of that product as a proxy for student competence. The primary threat is weakened inference, not merely increased misconduct (AERA et al., 2014; Kane, 2013; Messick, 1995; DEC, 2024).

Proposition 2. If widely available AI can complete an assessment to the required standard, the assessment was weak evidence for the intended learning claim

This is a validity statement, not a moral one. If the task can be satisfied through generic generation without demonstrating the targeted reasoning or judgment, then the interpretation and use argument is unstable under current conditions (Kane, 2013; Messick, 1995).

Proposition 3. Detection-first integrity strategies cannot restore validity at scale

Detection may have a limited diagnostic role, but it cannot be the primary mechanism for maintaining credibility because attribution is uncertain, incentives adapt quickly, and error costs create fairness and due process risk. Durable credibility requires redesign that strengthens the evidence base, not primarily stronger policing (TEQSA, 2025; Turnitin, 2025).

Proposition 4. Assessment vulnerability under AI conditions can be triaged using four likelihood subfactors

The likelihood that an assessment will fail as evidence under AI conditions is largely explained by: AI capability fit, task predictability, process observability, and assessment conditions. These subfactors provide a practical bridge between validity theory and portfolio-level governance (AERA et al., 2014; OECD, 2025b).

Proposition 5. High-stakes learning claims require high-observability evidence points

Programs can permit AI-assisted work in lower-impact tasks without collapsing credibility, but core outcomes and high-weight assessments require at least some evidence points where reasoning, decisions, and verification are observable enough to support defensible inference (DEC, 2025a; DEC, 2025b).

Proposition 6. Declared AI modes per assessment reduce ambiguity and improve governance

Explicitly declaring whether an assessment is Human only, Assist, AI–human interface, or Autopilot reduces the gap between formal policy and real practice. It also clarifies what the institution is actually assessing: production, or judgment and accountability with tools (HEPI, 2025; DEC, 2025b).

Proposition 7. Higher education's slow response is driven by predictable capacity, incentive, and psychology dynamics

Universities are structurally inclined toward legible policies and auditable procedures rather than design-heavy portfolio reform, producing decoupling between formal commitments and practice (Meyer & Rowan, 1977; Power, 1997; Bromley & Powell, 2012). Underreaction is also reinforced by psychic numbing and scope insensitivity, which blunt motivation when threats are diffuse and large-scale (Slovic, 2007; Kahneman, 2011).

Proposition 8. AI literacy is integrity infrastructure, and it must be embedded in disciplines

As assessment shifts toward accountable tool use and defensible judgment, AI literacy becomes necessary to meet assessment standards, not an optional "digital skills" add-on. This is especially important in constrained systems, where triage and reusable design patterns are the only scalable route to improved evidence and equity protection (European Commission & OECD, 2025; Molina & Medina, 2025; Reimers et al., 2026; DEC, 2024; DEC, 2025b).

9. Conclusion

Generative AI does not simply increase the probability of individual misconduct. It changes the conditions under which universities can make credible claims about learning from assessed work. When large parts of higher education rely on take-home products graded mainly on the final artifact, with limited visibility into process, the evidentiary value of that artifact declines as AI makes plausible production cheap and widely accessible (DEC, 2024; DEC, 2025b; OECD, 2025b). The resulting problem is system-level: a credibility gap between what assessment scores

are used to claim and what the assessment evidence can defensibly support. This gap is not resolved by retreating to invigilated written exams as a default remedy. Invigilation can strengthen identity assurance and increase control at the point of performance, but it does not remove AI from the assessment system because tool influence is readily displaced into preparation, rehearsal, and internalized scaffolding. In validity terms, supervision can tighten conditions without restoring the intended inference if capability fit remains high and tasks remain predictable (AERA et al., 2014; Kane, 2013; TEQSA, 2025).

The key point in this paper is validity-based rather than moralistic. Assessment is a claim about evidence. If widely available AI can complete an assessment to the required standard without the student demonstrating the targeted reasoning or judgment, the assessment was weak evidence for the intended learning claim. It may still have educational value, but it is unstable as a basis for high-stakes inference and downstream decisions unless redesigned (AERA et al., 2014; Kane, 2013; Messick, 1995).

FARABI is offered as a risk-based triage framework to respond proportionately. It shifts the unit of response from the suspected individual to the assessment portfolio. It asks where vulnerability is high (AI fit, predictability, observability, conditions) and where impact is high (weight, outcome centrality, transfer risk, compliance exposure), then concentrates redesign where improved evidence changes institutional inference most. It also forces explicit decisions about AI use modes, moving institutions away from the current reality of implicit, inconsistent practice (HEPI, 2025; TEQSA, 2025).

The framework is deliberately designed to avoid two failure modes that are already visible. The first is symbolic policy activity that leaves the underlying evidence system unchanged. The second is 'security theater' that leans on detection or restrictive controls that do not restore validity at scale and can create fairness and due process risk (Turnitin, 2025; TEQSA, 2025). This includes blanket reversion to supervised exams when it functions mainly as reassurance rather than as a redesign of the evidence chain (TEQSA, 2025). FARABI's "balanced integrity" stance treats proportionality, transparency, and equity as part of validity, not as optional ethics add-ons.

The paper also argues that FARABI functions as a Trojan horse. What looks like assessment repair is, in practice, an institutional design choice. If universities rebuild assessment to restore defensible evidence under AI conditions, they are committing to a model that privileges applied judgment, observable reasoning, accountable tool use, and discipline-embedded AI literacy. That shift is not a matter of preference or ideology. It follows from the logic of validity when output fluency is no longer scarce. The strategic implication is that universities are no longer deciding only how to assess. They are deciding what kind of institution they can credibly be at scale, and which learning claims they can still warrant (European Commission & OECD, 2025; Luckin, 2018).

Finally, the paper explains why higher education struggles to act even when the direction of travel is obvious. Reform is design-heavy and poorly rewarded by audit logics that privilege

legible procedures over improved inference, and institutional underreaction is reinforced by predictable psychological mechanisms, including psychic numbing and scope insensitivity (Slovic, 2007; Kahneman, 2011; Meyer & Rowan, 1977; Power, 1997). FARABI is intended to convert recognition into proportionate action without becoming a compliance factory.

This paper does not attempt to resolve the broader question of higher education's purpose if employability becomes a weaker organizing metric under AI conditions. The contribution here is narrower and urgent: a validity-grounded account of why assessment credibility is now a system problem, and a practical framework for rebuilding defensible evidence in ways that remain scalable, fair, and educationally meaningful.

Declaration of AI Use

During preparation of this manuscript, the author used ChatGPT to assist with grammar and readability and to check formatting consistency. The author takes full responsibility for the manuscript's claims, interpretations, and references.

References

AERA, APA & NCME. (2014). *Standards for Educational and Psychological Testing*. American Educational Research Association (AERA), American Psychological Association (APA) & National Council on Measurement in Education (NCME). https://www.testingstandards.net/uploads/7/6/6/4/76643089/standards_2014edition.pdf

Aoun, J. E. (2018). *Robot-proof: Higher education in the age of artificial intelligence* (paperback ed.). MIT Press. https://mitpress.mit.edu/9780262535977/robot-proof/

Biggs, J. (1996). Enhancing teaching through constructive alignment. *Higher Education, 32*(3), 347–364. https://doi.org/10.1007/BF00138871

Bretag, T., Harper, R., Burton, M., Ellis, C., Newton, P. M., Rozenberg, P., Saddiqui, S., & van Haeringen, K. (2018). Contract cheating: A survey of Australian university students. *Studies in Higher Education*. https://doi.org/10.1080/03075079.2018.1462788

Bromley, P., & Powell, W. W. (2012). From smoke and mirrors to walking the talk: Decoupling in the contemporary world. *The Academy of Management Annals, 6*(1), 483–530. https://doi.org/10.1080/19416520.2012.684462

DEC - Digital Education Council. (2025a). *Digital Education Council Global AI Faculty Survey 2025*. https://www.digitaleducationcouncil.com/post/digital-education-council-global-ai-faculty-survey

DEC - Digital Education Council (2025b) *The Next Era of Assessment A Global Review of AI in Assessment Design*. https://www.digitaleducationcouncil.com/post/the-next-era-of-assessment-a-global-review-of-ai-in-assessment-design.

DEC- Digital Education Council. (2024). *Digital Education Council Global AI Student Survey 2024*. https://www.digitaleducationcouncil.com/post/digital-education-council-global-ai-student-survey-2024

DiMaggio, P. J., & Powell, W. W. (1983). The iron cage revisited: Institutional isomorphism and collective rationality in organizational fields. *American Sociological Review, 48*(2), 147–160. https://doi.org/10.2307/2095101

European Commission. (2025). *AI literacy framework review draft* https://ailiteracyframework.org/

European Union. (2024). *Regulation (EU) 2024/1689 (Artificial Intelligence Act). Official Journal of the European Union.* https://eur-lex.europa.eu/eli/reg/2024/1689/oj

Freeman, J. (2025). *Student Generative AI Survey 2025* (HEPI/Kortext). Higher Education Policy Institute. https://www.hepi.ac.uk/reports/student-generative-ai-survey-2025/

Gibbs, G., & Simpson, C. (2004). Conditions under which assessment supports students' learning. *Learning and Teaching in Higher Education*, (1), 3–31. https://eprints.glos.ac.uk/3609/7/LATHE%20Contents%201.pdf

Gifford, R. (2011). The dragons of inaction: Psychological barriers that limit climate change mitigation and adaptation. *American Psychologist, 66*(4), 290–302. https://doi.org/10.1037/a0023566

Google. (2025). *2025 year in review*. https://blog.google/outreach-initiatives/education/google-for-education-year-in-review-2025/

Gulikers, J. T. M., Bastiaens, T. J., & Kirschner, P. A. (2007). Authentic assessment, student and teacher perceptions: The practical value of the five-dimensional framework. *Journal of Vocational Education and Training, 58*(3), 337–357. https://doi.org/10.1080/13636820600955443

Higher Education Policy Institute. (2025). *AI and the future of universities.* https://www.hepi.ac.uk/wp-content/uploads/2025/10/AI-and-the-Future-of-Universities.pdf

Kahneman, D. (2011). *Thinking, fast and slow*. Farrar, Straus and Giroux.

Kahneman, D., & Knetsch, J. L. (1992). Valuing public goods: The purchase of moral satisfaction. *Journal of Environmental Economics and Management, 22*(1), 57–70. https://doi.org/10.1016/0095-0696(92)90019-S

Kane, M. T. (2013). Validating the interpretations and uses of test scores. *Journal of Educational Measurement, 50*(1), 1–73. https://doi.org/10.1111/jedm.12000

Luckin, R. (2018). *Machine learning and human intelligence: The future of education for the 21st century*. UCL IOE Press. (Open PDF copy) https://discovery.ucl.ac.uk/id/eprint/10178695/

Messick, S. (1995). Validity of psychological assessment: Validation of inferences from persons' responses and performances as scientific inquiry into score meaning. *American Psychologist, 50*(9), 741–749. http://psycnet.apa.org/journals/amp/50/9/741/.

Meta. (2025, February 26). *New Meta for Education offering is now generally available.* https://about.fb.com/news/2025/02/new-meta-education-offering-generally-available/

Meyer, J. W., & Rowan, B. (1977). Institutionalized organizations: Formal structure as myth and ceremony. *American Journal of Sociology, 83*(2), 340–363. https://doi.org/10.1086/226550

Microsoft. (2024, September 18). *Bringing a new wave of Copilot innovation to education.* https://www.microsoft.com/en-us/education/blog/2024/09/bringing-a-new-wave-of-copilot-innovation-to-education/

Molina, E.; Medina, E. (2025). *AI Revolution in Higher Education: What you need to know*. Digital Innovations in Education; Brief No. 4. World Bank. https://doi.org/10.1596/43298

Newton, P. M. (2018). How common is commercial contract cheating in higher education and is it increasing? A systematic review. *Frontiers in Education, 3*, Article 67. https://doi.org/10.3389/feduc.2018.00067

European Commission & OECD. (2025). *Empowering Learners for the Age of AI: An AI Literacy Framework for Primary and Secondary Education (Review Draft)*. https://ailiteracyframework.org/wp-content/uploads/2025/05/AILitFramework_ReviewDraft.pdf

OECD. (2025a). *Trends Shaping Education 2025*. OECD Publishing. https://doi.org/10.1787/ee6587fd-en

OECD. (2025b). *Introducing the OECD AI capability indicators*. OECD Publishing. https://doi.org/10.1787/3c40105f-en

OpenAI. (2024, May 30). *Introducing ChatGPT Edu.* https://openai.com/blog/introducing-chatgpt-edu/

OpenAI. (2025a, February 25). *Estonia and OpenAI to bring ChatGPT to schools nationwide.* https://openai.com/index/estonia-schools-and-chatgpt/

OpenAI. (2025b, February 4). *OpenAI and the CSU system bring AI to 500,000 students and faculty.* https://openai.com/index/openai-and-the-csu-system/

Oxford University Press. (2025). *Teaching the AI-native generation: Empowering schools in the age of AI* (report). Oxford University Press. https://fdslive.oup.com/www.oup.com/oxed/secondary/Teaching_the_AI_Native_Generation.pdf

Power, M. (1997). *The audit society: Rituals of verification*. Oxford University Press. https://doi.org/10.1093/acprof:oso/9780198296034.001.0001

QAA. (2022). *Contracting to cheat in higher education: How to address contract cheating, the use of third-party services and essay mills* (3rd ed.). Quality Assurance Agency for Higher Education. https://www.qaa.ac.uk/docs/qaa/guidance/contracting-to-cheat-in-higher-education.pdf

Reimers, F., Azim, A., Palomo, M.-K., Thoney, C. (2026). *Artificial Intelligence and Education in the Global South*. Springer. https://doi.org/10.1007/978-3-032-11449-5.

Slovic, P. (2007). "If I look at the mass I will never act": Psychic numbing and genocide. *Judgment and Decision Making, 2*(2), 79–95. https://doi.org/10.1017/S1930297500000061

TEQSA. (2025, September 24). *Enacting assessment reform in a time of artificial intelligence*. https://www.teqsa.gov.au/guides-resources/resources/corporate-publications/enacting-assessment-reform-time-artificial-intelligence

Turnitin. (2025). *Crossroads: Navigating the intersection of AI and academia*. https://www.turnitin.com/whitepapers/ai-in-education.

Vlachopoulos, D., & Makri, A. (2024). A systematic literature review on authentic assessment in higher education: Best practices for the development of 21st century skills, and policy considerations. *Studies in Educational Evaluation, 83*, 101425. https://doi.org/10.1016/j.stueduc.2024.101425

Weber, E. U. (2006). Experience-based and description-based perceptions of long-term risk: Why global warming does not scare us (yet). *Climatic Change, 77*(1–2), 103–120. https://doi.org/10.1007/s10584-006-9060-3